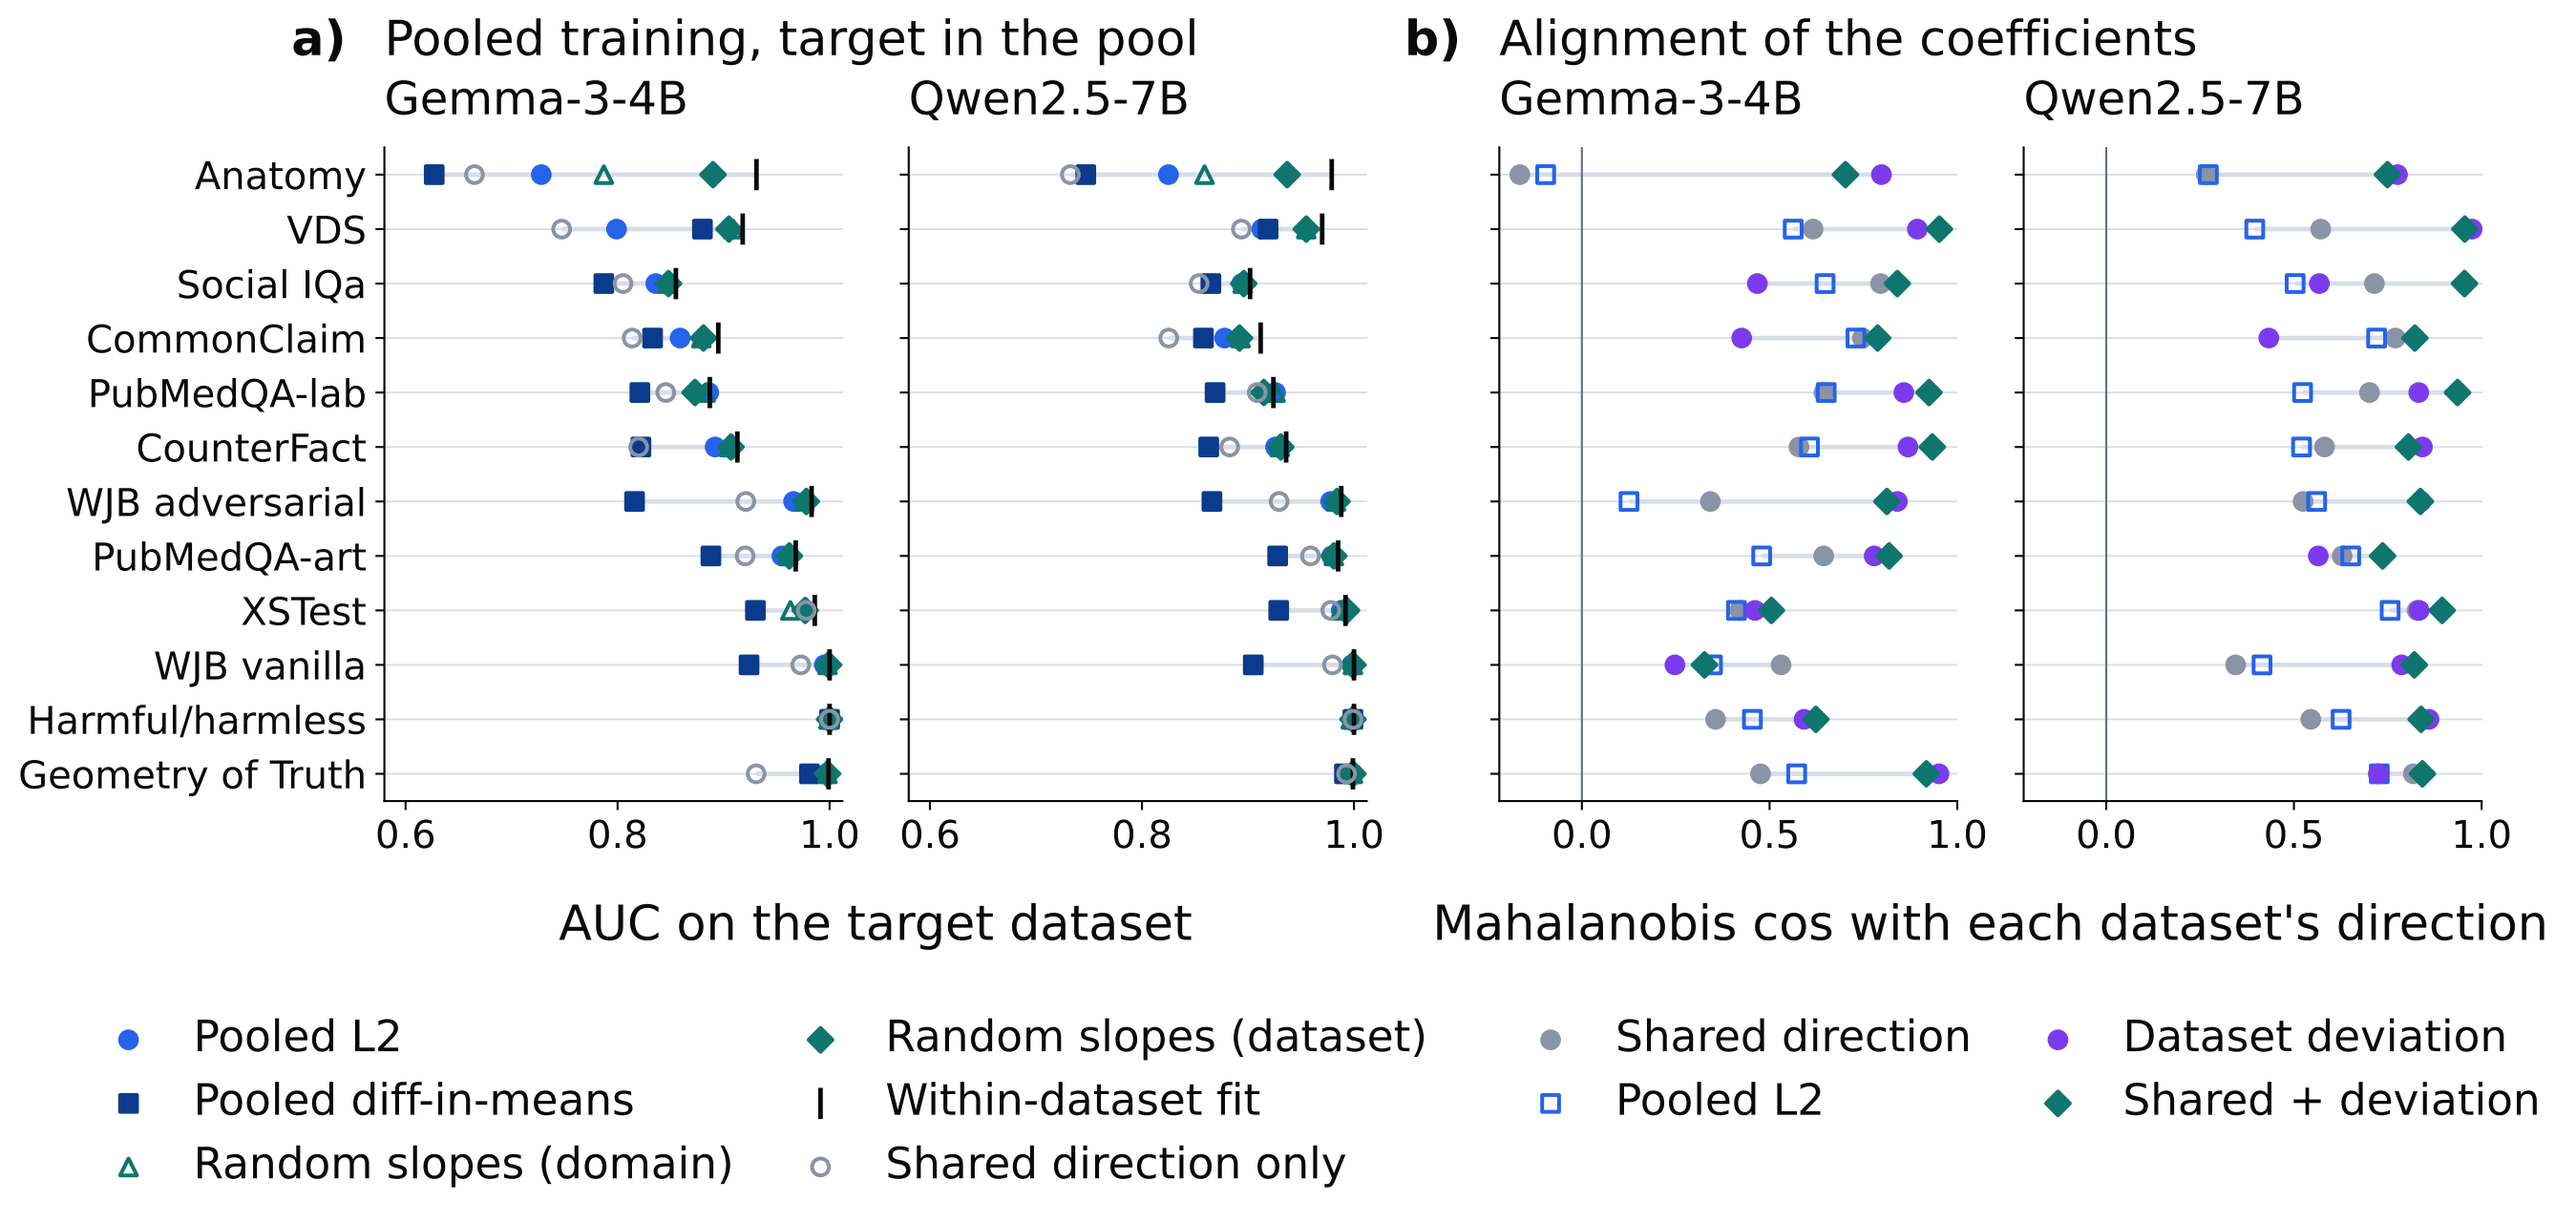
a) Pooled training, target in the pool
Gemma-3-4B
Qwen2.5-7B
Anatomy
VDS
Social IQa
CommonClaim
PubMedQA-lab
CounterFact
WJB adversarial
PubMedQA-art
XSTest
WJB vanilla
Harmful/harmless
Geometry of Truth
0.6
0.8
1.0
AUC on the target dataset
b) Alignment of the coefficients
Gemma-3-4B
Qwen2.5-7B
0.0
0.5
1.0
Mahalanobis cos with each dataset's direction
Pooled L2
Pooled diff-in-means
Random slopes (domain)
Random slopes (dataset)
Within-dataset fit
Shared direction only
Shared direction
Pooled L2
Dataset deviation
Shared + deviation